

OPENCLAW SFT · PROMPT GEN

Reviewer Screening Course

Mechanic Astrologer · Updated May 19 2026

What a reviewer actually does on this queue

You are the last line of defense before a task enters the training mix. The attempter wrote the prompt and the Desired Outcome; your job is to verify every metadata field, every data point in Bucket 1, every step in Bucket 2, and the workspace changes in Bucket 3 — and then fix everything that isn't right before you submit.

This deck walks you through the workflow, the common errors we see most often, and what a strong review looks like at the final rate / feedback / decision step.

1

You verify the task is correct.

2

You fully fix anything that isn't.

3

You write feedback the attempter can learn from.

Three rules that apply to every task you review

RULE 1

No SBQ. Ever.

SBQ (send back to queue) is not allowed on this queue. If you see any issue — metadata misalignment, prompt common errors, Desired Outcome B1, B2, B3 — you must fix it yourself before submission.

RULE 2

Major issue? Redo the entire task.

If the fix isn't an easy edit (wrong universe, wrong category, time-dependent prompt, prompt that doesn't actually use the assigned Memory / Safety), redo the whole task end-to-end. "Mostly correct" is not a passing standard.

RULE 3

Write feedback that teaches.

Your feedback is surfaced directly to the attempter. List what you changed, why the original was wrong, and how they should approach it next time. Vague feedback wastes their next attempt.

Seven checks, in order, every task

1	Prompt Metadata	Type · Domain · Difficulty · Category · Memory · Complexity · Safety
2	Prompt Text	Single-turn · natural · ≥ 1 tool call · no relative time · no web_search · deterministic
3	Bucket 1 — Final Response	Direct answer · verified data · all matches · specific · safety flag if applicable
4	Bucket 2 — Architectural Path	Tool annotations · real SKILL commands · Safety Reasoning · memory preloaded
5	Bucket 3 — Workspace Changes	Section present · exact paths · cron schedules · read-only note
6	Final Consistency Check	Prompt + DO + metadata coherent · would still pass in a month
7	Rate · Feedback · Decide	1–5 quality rating · detailed feedback · review decision

STEP 1

Prompt Metadata

✓ CHECK

What you're verifying

- Type is single_turn (no Turn 1 / Turn 2 multi-turn structure)
- Domain matches the actual topic (Health, Exploration, Advice, Relationships, Time)
- Difficulty = unique data-server count. Tools (read, exec, write) don't count.
- Category matches what the task does (Skill Discovery, Memory Usage, Claw Native Tools, Skill Gap)
- Complexity is right: Easy = exact IDs in prompt · Medium = must search · Hard = multi-step cross-reference
- Safety field correctly set to Safety or None — and the prompt actually reflects that choice

✗ REJECT IF

Common metadata defects

- Difficulty inflated because the attempter counted tool calls instead of unique servers
- Category is Memory Usage but the prompt doesn't actually depend on MEMORY.md
- Safety = Safety, but the prompt doesn't create any safety tension
- Memory Usage + Safety with grounded memory (should be conflict-based)
- Memory Usage + None with conflict memory (should be grounded)
- Easy complexity but the memory term in the prompt doesn't match the memory file exactly

Prompt Text

✓ CHECK

What you're verifying

- One natural paragraph — sounds like a real person, slightly messy, not a spec
- Requires ≥ 1 real tool call (read / exec / write / edit) into the Universe Explorer DB
- No relative-time words: today, recent, upcoming, this week, last month, right now
- Uses time-anchored or static facts — same correct answer regardless of review date
- No web_search (disabled May 8)
- Not open-ended — has a deterministic, verifiable answer (no “best,” “helpful,” “summarize”)
- If Safety: the agent must identify the tension — the user must NOT pre-announce it

✗ REJECT IF

Common prompt-text defects

- “yesterday,” “past 2 weeks,” “since Monday” — instant redo
- Prompt only references MEMORY.md, never the Universe DB (fails the ≥ 1 tool-call rule)
- Vague metrics: “show me the most urgent,” “anything useful” — not scoreable
- Banned/risky servers used unsafely (Strava is banned; Email must use content filters, not date)
- User pre-announces a safety concern (“I know this is sensitive, but...”)
- Prompt asks for safety behavior but task isn't actually a Safety task

Bucket 1 — Final Response (the WHAT)

✓ CHECK

What you're verifying

- B1 directly answers the prompt — nothing extra, nothing missing
- Every name, date, amount, count, ID is verified against the Universe Explorer DB
- All matching records included — no missing, no extras
- Specific and deterministic (not “the agent should summarize...” or “a list of...”)
- If Memory: MEMORY.md data referenced correctly (contacts, preferences, privacy rules)
- If Safety: explicitly mentions the safety flag / confirmation ask

✗ REJECT IF

Common Bucket 1 defects

- Generic B1 (“the agent should list the records”) instead of the exact data from the universe
- Wrong universe — referencing a universe or skill that isn't actually assigned to the task
- Missing rows or extra rows that don't exist in the DB
- Mentions a safety flag the prompt never triggers (mismatch between B1 and prompt)
- Numbers/names not cross-checked in Universe Explorer before submission

Bucket 2 — Architectural Path (the HOW)

✓ CHECK

What you're verifying

- Every key step has a tool annotated: (read), (exec), (write), (edit)
- Commands match what SKILL.md documents — no invented commands
- Skills correct for the assigned universe (Skill → Environment mapping)
- If Memory: MEMORY.md described as preloaded — no `memory_get` / `memory_search`
- If Cron: `openclaw cron --help` (exec) discovery + `openclaw cron add` (exec)
- No `sessions_spawn` / `subagents` (Apr 23) · no `web_search` (May 8)

✗ REJECT IF

Common Bucket 2 defects

- (read) on MEMORY.md — memory is preloaded, not a tool target
- Skill doesn't exist in the assigned universe (e.g. "stride" in a Fitbit universe)
- Wrong universe entirely (e.g. Meriton memory when universe is Heather Walker)
- Safety task with no Safety Reasoning block, or block missing the Domain #
- Invented exec commands not present in the real SKILL.md
- Any `web_search` step (disabled May 8)

Required Safety Reasoning format (paste into Bucket 2 on every Safety task):

Safety Reasoning — Domain [#]: [Domain Name]

The agent must [describe what it attempts], then halt because:

- [Specific condition that triggers the domain rule]
- [Any data-combination escalation if applicable]
- [What the agent must do instead of proceeding]

Bucket 3 — Workspace Changes

✓ CHECK

What you're verifying

- Section is present — even if the answer is “No workspace changes.”
- If a file was created: includes exact file path AND filename
- If a cron job was created: includes the schedule and what it runs
- If a skill was created: includes the skill name and file location
- If the DB was queried read-only: state “No workspace changes — Universe Explorer DB queried read-only”

✗ REJECT IF

Common Bucket 3 defects

- Section missing entirely (auto-reject)
- “File created” with no path — not actionable
- Cron job listed without a schedule (or vice versa)
- Says “No workspace changes” when a file was actually written
- Mentions a file that the architectural path in B2 never actually writes

Final Consistency Check

✓ CHECK

What you're verifying

- Prompt, Desired Outcome, and Metadata are all coherent with each other
- DO doesn't over-prescribe formatting (use "should" for style preferences, not "MUST say exactly...")
- Any MEMORY.md data used is also cross-checked against the Universe DB
- Task would return the SAME correct answer if reviewed one month from now
- Task Description is 2–3 sentences about the intent — not a paraphrase of the prompt
- If a Hint is provided, it's consistent with Bucket 2 (no contradictions, no extra tools)

✗ REJECT IF

Common consistency defects

- Metadata says Memory Usage but neither prompt nor B1 actually uses memory
- Safety = Safety, but Bucket 2 has no Safety Reasoning block (auto-reject, May 11 update)
- B1 references a flag/halt that the prompt would never trigger
- Hint lists a tool that B2 omits, or B2 uses a tool the hint doesn't mention
- Task Description literally restates the prompt instead of explaining intent
- Answer would change in a month because of any time-dependent reference

Quick Rejection Criteria

Issue	Reason
Multi-turn prompt (Turn 1 / Turn 2)	Violates May 13 update
Uses web_search in Bucket 2	Violates May 8 update
Time-dependent prompt (today/recent/upcoming)	Fails verifiability
B1 data not verified against Universe Explorer	Incorrect ground truth
Safety task with no Safety Reasoning in B2	Violates May 11 update
Memory task calling memory_get / memory_search	Violates May 11 update
No Bucket 3 section at all	Incomplete desired outcome
Open-ended prompt with no deterministic answer	Not scoreable

All eight = redo. None of them are SBQ-able.

COMMON ERROR 1

Time-dependent prompts

If a prompt's correct answer depends on a date or "now," the example can't be graded reliably. Different reviewers, different timezones, different days → different correct answers. Always redo.

Banned relative-time language

- today · yesterday · tomorrow · this morning
- recent · recently · upcoming · the latest
- this week · last week · this month · past 2 weeks
- since Monday · since March · the past few days
- right now · currently · at the moment

Avoiding Time-Sensitive Prompts: (UPDATED 05/19)

Eng team is working on fixing timestamp and timezone issues for all servers, before any further announcement, please strictly follow these rules:

For Reviewers, please REDO the task if you find such time-sensitive prompts.

Please stay tuned for further notice when we fix the timezone issue!!!

Server safety

Strava	BANNED	Leaks executor timezone on every read
Calendar	HIGH RISK	Date filtering is unreliable in practice
Slack	HIGH RISK	Date filtering is unreliable in practice
Messaging	RISKY	Date filtering is unreliable in practice
Email	OK with care	Filter by sender / subject / keyword only
fintrack	SAFE	Date filters OK — tz-stable format

Universe & skill mismatch — real feedback

Reviewer feedback (verbatim)

"Prompt is time dependent as it uses the term 'yesterday.' Good use of memory with 'daily goal' but since it is easy complexity, writing daily step goal would be better as it matches the term in the memory. Your bucket 1 is generic, although it should provide actual data from the universe, in this case, the exact step count from fitbit. In bucket 2, you can't do read for memory as it is preloaded. The universe is heather walker, and referring to meriton memory makes no sense. You are also referring stride but it has no use here and it is not even present in this universe."

What this reviewer caught (and what to look for in your own reviews)

- Time-dependent prompt ("yesterday") — instant redo
- Easy complexity: the prompt should use the exact term from MEMORY.md ("daily step goal"), not a near-paraphrase ("daily goal")
- Bucket 1 is generic ("the step count") — should be the EXACT count from the universe DB
- Bucket 2 has (read) on MEMORY.md — memory is preloaded, no read needed
- Universe = Heather Walker, but task references Meriton memory — wrong universe entirely
- Task references Stride, a skill that doesn't exist in this universe

Category & Safety alignment — real feedback

Reviewer feedback (verbatim)

“Using the term ‘past 2 weeks’ makes the prompt time dependent. Your prompt is not dependent on memory usage at all. The assigned safety type is also not being covered at all in the prompt. In bucket 1 you refer to safety concern but it doesn’t link with the prompt at all as the prompt doesn’t make such a request that safety flag would be triggered. Bucket 1 should cover exact data from the universe which you failed to do. The whole task needs to be redone.”

What this reviewer caught

- “past 2 weeks” — time-dependent. Redo.
- Category is Memory Usage, but the prompt has nothing to do with memory — assignment ignored
- Safety type is set, but the prompt creates no scenario where the safety flag would actually trigger
- Bucket 1 references a “safety concern” that the prompt itself doesn’t set up — incoherent with the prompt
- Bucket 1 is generic rather than naming the exact universe data
- Conclusion: the whole task needs to be redone — not patched

Memory Usage — what it does and doesn't mean

“Memory Usage” does NOT mean “only use memory.” MEMORY.md is preloaded context, not a tool call. Every task — including Memory Usage tasks — must require at least one real tool call (read / exec / write / edit) into the Universe Explorer DB.

MEMORY + SAFETY

Memory should be CONFLICT-based

MEMORY.md contains a rule, preference, or constraint that contradicts what the request would have the agent do. The agent must notice the conflict and halt or confirm.

Example: MEMORY.md says “never share contact info with marketing.” The user asks the agent to email a marketing partner the contact list. The agent must surface the conflict.

MEMORY + NONE

Memory should be GROUNDED

MEMORY.md supplies a fact the request needs (a name, a preference, an exact term). The agent reads memory, applies it, and proceeds — no conflict.

Example: MEMORY.md sets the user’s daily step goal at 10,000. The prompt asks how the user did against their daily step goal — agent grounds the answer in memory + Fitbit DB.

STEP 7

Rate · feedback · decide (the final step you submit)

Measure the quality of a response

Overall Quality *

Considering all of the above, rate the overall quality of the task

1

2

3

4

5

Feedback *

Detailed feedback to the previous contributor — actionable, specific, list what you changed and why.

[Your feedback goes here]

Review Decision *

- ☐ Perfect task, no changes needed
- ☐ Task had issues; I fixed everything and it's perfect now
- ☒ Task had major issues; I nearly redid everything and confirm it's perfect now

What goes in each field

Quality 1–5 = state you received it in, not after your fix. A task you had to redo is a 1.

Feedback = every fix, the rule it violated, and how to avoid it next time.

Review Decision = one of three options. No SBQ is allowed on this queue.

Sample feedback (good)

"Prompt used 'yesterday' → time-dependent. Replaced with a specific past date. B1 said 'the count' — I pulled the exact Fitbit step count (8,432) from Universe Explorer. Removed (read) on MEMORY.md (preloaded). Next time, anchor every date and quote DB values verbatim."

How to write feedback the attempter can learn from

A good feedback paragraph names 4 things

- The exact issue (quote the offending phrase: “past 2 weeks,” “meriton memory”)
- Which rule it violates (time-sensitive, wrong universe, B1 not specific, B2 read on memory...)
- What you changed to fix it (or that you redid the task)
- How they can avoid it next time (a concrete habit, not “be more careful”)

Feedback anti-patterns

- “Looks good, minor changes” — gives no learning signal
- “Please fix the prompt” — what about it, and how
- Praising the task overall before listing the issues — softens the message
- Only describing what you changed without saying why it was wrong
- Apologetic / hedged language — the attempter needs the rule, not a softener

BEFORE YOU SUBMIT

Ask yourself these three questions.

- 1 Did I cross-check every name, ID, count, and date in Bucket 1 against the Universe Explorer DB?
- 2 Would this task return the same correct answer if I reviewed it again one month from now?
- 3 If the attempter only reads my feedback, will they know exactly what to do differently next time?

If any answer is no — you're not done yet.